

Experiment -1

- 1. For the given data, analyze the missing values and comment on imputing or dropping the missing values. Explain how the missing values will affect the model estimation and its goodness of fit.**
- 2. Influence of variance in large data set with many variables. Higher the variance of the variable will contribute the higher importance to the data set. Find the variance of method to reduce the dimension of the data set. Give justification of output.**
- 3. Dimensionality reduction using PCA (Principal Component Analysis) and Feature Selection through OMP (Orthogonal Matching Pursuit)**

Aim: To deal with a dataset having missing values and analyze how the missing values will affect the model estimation and its goodness of fit.

Description:

- Create a dummy dataset.
- Use different aspects of missing data imputation
- Find out variances of all the variables and by including and excluding different variables, observe the effects.
- Apply PCA (Principal Component Analysis) and Feature Selection through OMP (Orthogonal Matching Pursuit).

Algorithm:

- Missing data imputation
- Calculate variances of all the variables.
- Apply PCA (Principal Component Analysis) and Feature Selection through OMP (Orthogonal Matching Pursuit) on the given dataset.

Program:

```
####test for missing values
```

```
#vector with missing values
```

```
x<-c(1:4,NA,6:7,NA)
```

```
x
```

```
#test the missing values in x vector
```

```
is.na(x)
```

```
#data frame with missing values
```

```
df<-data.frame(col1=c(10:12,NA),  
               col2=c("I",NA,"am","fine"),  
               col3=c(TRUE,FALSE,TRUE,FALSE),  
               col4=c(1.2,2.3,NA,4.3),  
               stringsAsFactors=FALSE)
```

```
df
```

```
#IDENTIFY THE MISSING VALUES IN THE DATA FRAME
```

```
is.na(df)
```

```
#identify the location of NA's in the vector
```

```
which(is.na(x))
```

```
#identify the location of NA's in the data frame
```

```
which(is.na(df))
```

```
#display the sum of NA's in vector
```

```
sum(is.na(x))
```

```
#display the sum of NA's in the data frame
```

```
sum(is.na(df))
```

```
#display the count of NA's in the each column of data frame
```

```
colSums(is.na(df))
```

```
#in two ways missing values are coded i.e., NA and 99
```

```
#recode missing values with mean
```

```
#before recode x is
```

```
x
```

```
mean(x)
```

```
mean(x,na.rm=TRUE)
```

```
x[is.na(x)]<-mean(x,na.rm=TRUE)
```

```
#after recode x is
```

```
x
```

```
round(x,2)
```

```
# data frame with missing values coded as 99
```

```
df1<-data.frame(col1=c(1:3,99),col2=c(2.8,4.7,99,4.2))
```

```
df1
```

```
#change 99 to NA
```

```
df1[df1==99]<-NA
```

```
df1
```

```
#Exclude missing values
```

```
x1<-c(1:3,NA,4:5)
```

```
x1
```

```
#display mean of x1
```

```
mean(x1)
```

```
#mean of x1 after excluding NA
```

```
mean(x1,na.rm=TRUE)
```

```
#DATA FRAME WITH MISSING VALUES
```

```
df
```

```
#list of complete rows of the data frame
```

```
complete.cases(df)
```

```
#subset with complete.cases to get complete cases
```

```
df[complete.cases(df),]
```

```
#subset with ! to get incomplete cases
```

```
df[!complete.cases(df),]
```

```
#short hand way to get the complete cases of data frame by omitting
```

```
na.omit(df)
```

Output:

```
> x
```

```
[1] 1 2 3 4 NA 6 7 NA
```

```
> is.na(x)
```

```
[1] FALSE FALSE FALSE FALSE TRUE FALSE FALSE TRUE
```

```
#data frame with missing values
```

```
> df
```

	col1	col2	col3	col4
1	10	I	TRUE	1.2
2	11	<NA>	FALSE	2.3
3	12	am	TRUE	NA
4	NA	fine	FALSE	4.3

```
> is.na(df)
```

	col1	col2	col3	col4
[1,]	FALSE	FALSE	FALSE	FALSE
[2,]	FALSE	TRUE	FALSE	FALSE
[3,]	FALSE	FALSE	FALSE	TRUE
[4,]	TRUE	FALSE	FALSE	FALSE

```
> which(is.na(x))
```

```
[1] 5 8
```

```
> sum(is.na(x))
```

```
[1] 2
```

```
> sum(is.na(df))
```

```
[1] 3
```

```
> colSums(is.na(df))
```

col1	col2	col3	col4
1	1	0	1

#in two ways missing values are coded i.e., NA and 99

#recode missing values with mean

#before recode x is

```
> x
```

```
[1] 1 2 3 4 NA 6 7 NA
```

```
> mean(x)
```

```
[1] NA
```

```
> mean(x,na.rm=TRUE)
```

```
[1] 3.833333
```

```
> x[is.na(x)]<-mean(x,na.rm=TRUE)
```

```
> #after recode x is
```

```
> x
```

```
[1] 1.000000 2.000000 3.000000 4.000000 3.833333 6.000000 7.000000  
[8] 3.833333
```

```
> round(x,2)
```

```
[1] 1.00 2.00 3.00 4.00 3.83 6.00 7.00 3.83
```

```
> # data frame with missing values coded as 99
```

```
> df1<-data.frame(col1=c(1:3,99),col2=c(2.8,4.7,99,4.2))
```

```
> df1
```

```
  col1 col2  
1     1  2.8  
2     2  4.7  
3     3 99.0  
4    99  4.2
```

```
> #change 99 to NA
```

```
> df1[df1==99]<-NA
```

```
> df1
```

```
      col1 col2
1       1  2.8
2       2  4.7
3       3   NA
4      NA  4.2
```

```
> #Exclude missing values
```

```
> x1<-c(1:3,NA,4:5)
```

```
> x1
```

```
[1]  1  2  3 NA  4  5
```

```
> #display mean of x1
```

```
> mean(x1)
```

```
[1] NA
```

```
> #mean of x1 after excluding NA
```

```
> mean(x1,na.rm=TRUE)
```

```
[1] 3
```

```
> #DATA FRAME WITH MISSING VALUES
```

```
> df
```

```
      col1 col2  col3 col4
1      10    I  TRUE  1.2
2      11 <NA> FALSE  2.3
3      12   am  TRUE   NA
4      NA fine FALSE  4.3
```

```
> #list of complete rows of the data frame
```

```
> complete.cases(df)
```

```
[1]  TRUE FALSE FALSE FALSE
```

```

> #subset with complete.cases to get complete cases

> df[complete.cases(df),]

  col1 col2 col3 col4
1   10    I TRUE  1.2
> #subset with ! to get incomplete cases

> df[!complete.cases(df),]

  col1 col2  col3 col4
2   11 <NA> FALSE  2.3
3   12   am  TRUE   NA
4    NA fine FALSE  4.3

> #short hand way to get the complete cases of data frame by omitting

> na.omit(df)

  col1 col2 col3 col4
1   10    I TRUE  1.2

```

Experiment -1 (2)

Influence of variance in large data set with many variables. Higher the variance of the variable will contribute the higher importance to the data set. Find the variance of method to reduce the dimension of the data set. Give justification of output.

Aim: To find out the influence of variance in large data set with many variables.

Description:

- Variance has a huge impact in many different aspects of your life.
- Variance measures how distant from the mean random values are in a data set.
- A set of data with low variance (relative) is dominated at the mean, and a set of high variance is spread out and deviates significantly from the mean.

- A high variance curve will be flat relative to a low variance curve.

Algorithm:

- Variance is mathematically defined as the expected value of the squared deviation from the mean. It can also be calculated from a data set by using the following equation:

$$\sigma^2 = \frac{\sum (\chi - \mu)^2}{N}$$

```
Corrmatrix <- structure(c(1, 0.82, 0.54, 0.36, 0.85, 0.82, 1, 0.01, 0.74, 0.36,
                          0.54, 0.01, 1, 0.65, 0.91, 0.36, 0.74, 0.65, 1, 0.36,
                          0.85, 0.36, 0.91, 0.36, 1),
                        .Dim = c(5L, 5L))
```

Corrmatrix

```
library(caret)
```

```
findCorrelation(Corrmatrix, cutoff = .6, verbose = TRUE, names = F)
```

Output:

```
> Corrmatrix
```

```
      [,1] [,2] [,3] [,4] [,5]
[1,] 1.00 0.82 0.54 0.36 0.85
[2,] 0.82 1.00 0.01 0.74 0.36
[3,] 0.54 0.01 1.00 0.65 0.91
[4,] 0.36 0.74 0.65 1.00 0.36
[5,] 0.85 0.36 0.91 0.36 1.00
```

```
> findCorrelation(Corrmatrix, cutoff = .6, verbose = TRUE, names = F)
```

Compare row 1 and column 5 with corr 0.85

Means: 0.642 vs 0.545 so flagging column 1
Compare row 5 and column 3 with corr 0.91
Means: 0.543 vs 0.499 so flagging column 5
Compare row 3 and column 4 with corr 0.65
Means: 0.33 vs 0.352 so flagging column 4

All correlations ≤ 0.6

[1] 1 5 4

Justification:

If two variables having high covariance, one of them may be removed from the data since both of them will be equally contributing individually.

Experiment -1 (3)

Dimensionality reduction using PCA (Principal Component Analysis) and Feature Selection through OMP (Orthogonal Matching Pursuit).

Aim: Performing the dimensionality reduction operation using PCA (Principal Component Analysis) and Feature Selection through OMP (Orthogonal Matching Pursuit).

Implementing PCA:

Algorithm:

1. Load factoextra for visualization
2. Compute PCA
3. Visualize eigenvalues (scree plot). Show the percentage of variances explained by each principal component.
4. Graph of individuals. Individuals with a similar profile are grouped together.
5. Graph of variables. Positive correlated variables point to the same side of the plot.
Negative correlated variables point to opposite sides of the graph.
6. Biplot of individuals and variables

```
library("factoextra")
```

```
data(decathlon2)
```

```

decathlon2.active <- decathlon2[1:23, 1:10]
head(decathlon2.active[, 1:6])

## Compute PCA
res.pca <- prcomp(decathlon2.active, scale = TRUE)

## Visualize eigenvalues (scree plot). Show the percentage of variances explained by each
principal component.
fviz_eig(res.pca)

# Graph of individuals. Individuals with a similar profile are grouped together.
Fviz_pca_ind(res.pca,
  col.ind = "cos2", # Color by the quality of representation
  gradient.cols = c("#00AFBB", "#E7B800", "#FC4E07"),
  repel = TRUE    # Avoid text overlapping
)

## Graph of variables. Positive correlated variables point to the same side of the plot. Negative
correlated variables point to opposite sides of the graph.
fviz_pca_var(res.pca,
  col.var = "contrib", # Color by contributions to the PC
  gradient.cols = c("#00AFBB", "#E7B800", "#FC4E07"),
  repel = TRUE    # Avoid text overlapping
)

## Biplot of individuals and variables
fviz_pca_biplot(res.pca, repel = TRUE,
  col.var = "#2E9FDF", # Variables color
  col.ind = "#696969" # Individuals color
)

```


Result:

```
> head(decathlon2.active[, 1:6])
```


	X100m	Long.jump	Shot.put	High.jump	X400m	X110m.hurdle
SEBRLE	11.04	7.58	14.83	2.07	49.81	14.69
CLAY	10.76	7.40	14.26	1.86	49.37	14.05
BERNARD	11.02	7.23	14.25	1.92	48.93	14.99
YURKOV	11.34	7.09	15.19	2.10	50.42	15.31
ZSIVOCZKY	11.13	7.30	13.48	2.01	48.62	14.17
McMULLEN	10.83	7.31	13.76	2.13	49.91	14.38

```
> ## Compute PCA
```


```
> res.pca <- prcomp(decathlon2.active, scale = TRUE)
```



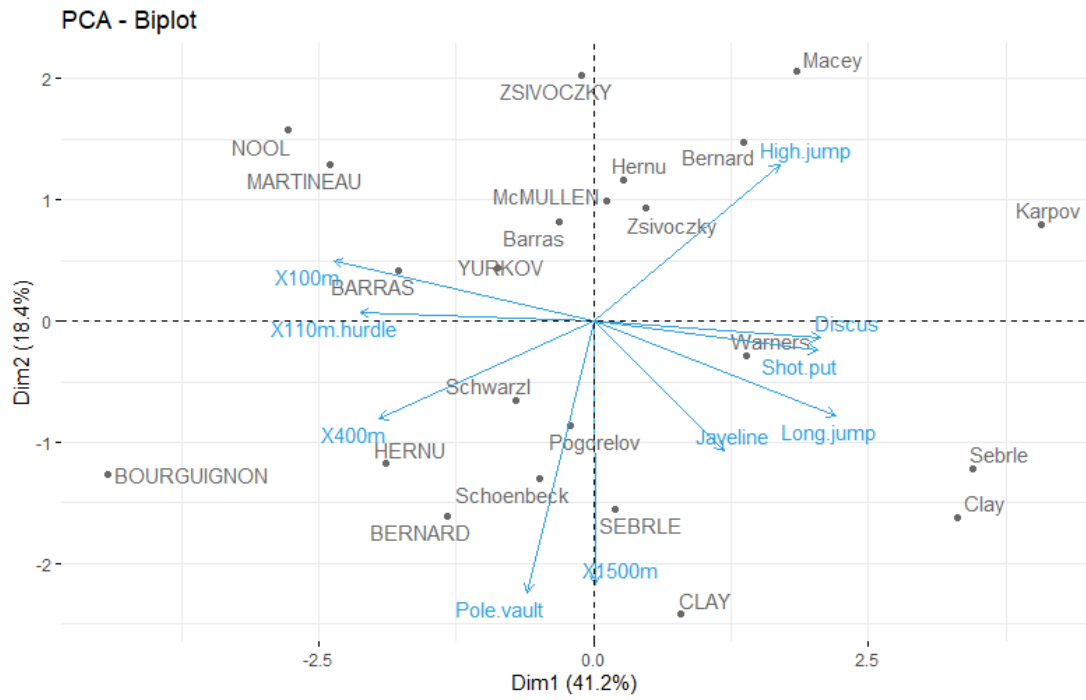
```
> ## Graph of individuals. Individuals with a similar profile are grouped together.
```



Graph of variables. Positive correlated variables point to the same side of the plot. Negative correlated variables point to opposite sides of the graph.



Biplot of individuals and variables



Implementing Feature Selection through OMP (Orthogonal Matching Pursuit)

